

○

ML HACKATHON

Analysis_Report: PESU_EC_030_032_035_045

CONTRIBUTORS:

Devanshi Shukla (PES2UG23AM030)

Dhriti Rai (PES2UG23AM032)

Farhaan Arshad (PES2UG23AM035)

Ishaan Tandon (PES2UG23AM045)

1) Key Observations

- Hardest parts

- **State representation design.** Encoding a variable-length masked word into a fixed-size numeric vector required a careful choice.
- **Reward shaping.** Getting a sparse but meaningful reward signal in Hangman was a little tough. We needed to guess correctly while also penalizing wasted guesses and repeats.
- **Combining probabilistic model with RL.** Using the HMM to provide prior letter probabilities helped guide the agent's action preferences, but integrating those probabilities properly into the observation vector and into the mask logic required iteration.

- Insights gained

- A simple char-level HMM over the corpus dataset gives a useful prior that directs early exploration toward letters likely given the pattern.
 - Small positive reward per correct letter and a larger terminal reward for finishing the word accelerates learning more than a purely terminal-only reward.
 - Masking invalid actions and using a target network stabilizes learning.
 - Curriculum or more corpus coverage would likely improve generalization.
-

◦

2) Strategies (design choices & why)

HMM design

- **Class & training:** We implemented a character-level HMM. Training is counting first-letter frequencies and letter-to-letter transitions (`start_counts`, `trans_counts`), with Laplace smoothing. Log-probabilities (`log_start`, `log_trans`) are used for numerical stability.
- **Scoring & letter distribution:** For a given pattern (masked word), `letter_distribution(pattern, guessed, words)`:
 - Filters corpus words by pattern length and matching observed letters.
 - Scores each candidate word with `logprob_word` (sum of start + transition log-probs).
 - Converts log-probs to normalized weights and accumulates probabilities for letters at blank positions.
 - Zeros out probabilities for already-guessed letters and normalizes the final letter probability vector.
- A char-level HMM is compact, fast, and gives a statistically grounded prior for letter likelihoods conditioned on the masked pattern. It complements RL by providing informed priors early in training and for unseen states.

RL (environment & agent)

- **Environment:** `HangmanEnv` models the game. `step(letter)`:
 - If letter already guessed → reward = -2.
 - If letter in word → reveal positions, reward = +1.
 - If letter not in word → reward = -1 and increment wrong.
 - If the word becomes fully revealed → reward += 10 and `done` = True (win).
 - If wrong $\geq \text{max_wrong}$ → reward -= 5 and `done` = True (lose).
- **Observation / State (`obs_to_vec`)**
 - pattern encoded as one-hot-like per position: `MAX_LEN * 27` (each position encoded across 27 dims — 26 letters + blank). This gives a positional encoding of which letter is present or blank.

- - guessed vector of length 26 (binary hot-vector for which letters have been guessed).
 - hmm_probs vector of length 26 (the HMM's letter probabilities for the current pattern).
 - lives normalized scalar (lives/6).
 - Total STATE_DIM = MAX_LEN * 27 + 26 + 26 + 1.
 - **Action space:** 26 discrete actions (letters a–z). Actions masked via mask_actions(obs) which sets guessed letters' mask to 0. During learning the Q-values for masked actions are penalized by adding a large negative term (-1e6) so they are never chosen.
 - **Agent architecture:** DQN (PyTorch) feed-forward network:
 - DQN(s_dim, a_dim, hidden=256) with layers:
 - Linear(s_dim → 256) → ReLU
 - Linear(256 → 128) → ReLU
 - Linear(128 → a_dim)
 - policy_net and target_net used, with MSE loss and Adam optimizer (lr 1e-4).
 - **Replay & training details**
 - ReplayBuffer capacity = 50,000.
 - learn_start = 2000, learn_every = 4.
 - Batch size = 64.
 - Gamma = 0.99.
-

○

3) Exploration

- **Exponential decay:** Smoothly reduces randomness so the agent has broad experience early (important to populate replay buffer with varied transitions) and then focuses on exploiting learned Q-values.
 - **Action masking:** Masking guessed letters both in the action-selection step and in the training target (big negative Q for masked actions) reduces wasted exploration and makes exploration more meaningful (explore untried letters rather than repeats).
-

4) Future Improvements

- **More careful reward shaping & ablation**
 - Experiment with differing magnitudes for per-letter reward vs terminal reward (e.g., smaller per-letter reward but larger terminal win reward), or use a shaped reward proportional to letters revealed at the guess.
- **Better language model**
 - Replace or augment HMM with a word-level language model or an n-gram / masked language model to get stronger prior distributions over candidate words.
- **State representation upgrades**
 - Use embeddings for letters/positions (learnable) rather than fixed 27-dim position one-hot, or use a small recurrent encoder (LSTM) over the pattern to compress variable-length pattern information more effectively.
- **Policy improvements**

Try actor-critic (A2C) or PPO style policy gradient methods which sometimes handle sparse/episodic rewards better than DQN.

 - Alternatively, treat the problem as ranking candidate letters (use cross-entropy / supervised pretraining from corpus-derived heuristics).
- **Data & training**
 - Augment corpus (more words) and create curriculum training (start with short words, then gradually increase length).
 - Increase episodes and tune hyperparameters (batch, lr, hidden sizes).
- **Evaluation & metrics**

- - Log and compare metrics: success rate, avg wrong guesses, avg repeated guesses, distribution of wins/loses by word length.
 - Evaluate generalization to out-of-corpus words and to different word lengths.

- **Safety & robustness**

- Ensure masking is airtight; add checks so that the policy is never able to select masked actions at eval time.

- **Architecture tweaks**

- Try deeper/wider nets, or small attention layers to focus on pattern positions most informative for current HMM probs.